

# Rendering: Next Event Estimation

Bernhard Kerbl



Research Division of Computer Graphics  
Institute of Visual Computing & Human-Centered Technology  
TU Wien, Austria



- At this point, we should mention that we won't be needing uniform hemisphere sampling anymore...
- We can replace it for most purposes with BRDF sampling, that is, importance sampling the hemisphere based on the material
  - E.g., for diffuse materials, cosine-weighted hemisphere sampling
  - We will see solutions for other materials soon!
- BRDF sampling usually improves quality in most cases (how much?)



# Diffuse Path Tracing: Uniform vs BRDF Sampling



- Cosine-weighted hemisphere sampling “works”... can we do better?




# How to Get Rid of All that Noise?

- Higher-dimensional path tracing is particularly prone to noise
- How can we reduce noise in our renderings?
- Common suggestions when looking for ways to reduce noise:
  - Use more samples (brute force, often takes too much time)
  - Use importance sampling (already applied)
  - Use today's technique, **Next Event Estimation (NEE)**
    - Based on something we saw before: light source sampling



# Revisiting the Math Behind Light Source Sampling

## Direct lighting (soft shadows) (usable for rendering)



# Revisiting the Math Behind Light Source Sampling

## Direct lighting (soft shadows) (the same, but more explicit)

$$L_e^{[l]}(x) = \int_{S_l} f_r(x, y \rightarrow v) L_e^{[l]}(y) V(x, y) \left( \frac{y - x}{|y - x|} \cdot n_x \right) \frac{\left( \frac{x - y}{|x - y|} \cdot n_y \right)}{|x - y|^2} dA_y$$

Light going in direction  $v$

Material, modelled by the BRDF

emitter  $\cos(\theta)$

receiver  $\cos(\theta)$

distance

Light intensity at position  $y$  on the surface

visibility (new, ray tracing)





# Direct Lighting with Light Source Sampling

```
function Li(v_inv)
    x = scene.trace(v_inv)       $f_r(x, \omega \rightarrow v) L_e^{[l]} V(x, y) \cos \theta \frac{(-\omega \cdot n_y)}{|x - y|^2} A^{[l]}$ 
    f = x.emit
    y, area = light_surface_uniform_world()
    omega = (y-x).normalised()
    r = make_ray(x, omega)
    v = 0
    if (scene.trace(r) == y)
        v = 1
    P = dot(y.normal, -omega) / dot(y-x, y-x)
    f = x.alb/pi * y.emit * v * dot(x.normal, omega) * P * area
return r
```





# Direct Light: BRDF Sampling vs Light Source Sampling



- In both cases, you want to find out from how many directions on the hemisphere a point does receive how much light
- With uniform or BRDF sampling, pick one direction for each sample, and pretend that this direction speaks for the entire hemisphere
  - But if you collect many of these estimates and average, you converge
- With light source sampling, pick a point  $y$  on the light, and pretend the evaluation (can  $y$  illuminate  $x$ ?) holds for its entire surface area
  - But if you collect many of these estimates and average, you converge



- Both BRDF and Light Source Sampling perform importance sampling
- They selectively put samples at opportune places on the domain
  - BRDF: assuming uniform lighting, which directions contribute most?
  - Light source: knowing light locations, which directions may hit them?




# Exploiting Light Source Sampling for Path Tracing?


- Can light source sampling help us with path tracing?
  - Based on projecting all known light sources onto the hemisphere
  - Every surface in the scene is a potential source of (indirect) light!
  - If we treat every surface as a potential light source, we are back to randomly sampling the full hemisphere...
- Idea: follow each ray via multiple indirect bounces, but at each bounce, compute the direct lighting from light source surfaces!
  - Detected light at each bounce is no longer dependent on coincidence
  - This is what we refer to as ***Next Event Estimation***



- Builds on light source sampling. Think: where can light come from?




- Builds on light source sampling. Think: where can light come from?




- We can map out the full hemisphere and distinguish direct/indirect

indirect


direct




- At each bounce, use light source sampling to get direct illumination
- Sample the BRDF to create direction for collecting indirect light



- At each bounce, use light source sampling to get direct illumination
- Sample the BRDF to create direction for collecting indirect light




- At each bounce, use light source sampling to get direct illumination
- Sample the BRDF to create direction for collecting indirect light



- Light source sampling for direct light

+

BRDF sampling for indirect light



- Add them together to cover the hemisphere

- Light source sampling to project light source onto hemisphere
- Importance sampling of the hemisphere via the BRDF to generate next direction to collect potential indirect light from next hit point




# First Attempt at Next Event Estimation

```
function Li(v_inv, D)
...
f = x.emit
...
// Apply Russian Roulette at some point
...
direct = <direct lighting with light source sampling>
...
indirect = <indirect light (recursive) with BRDF sampling>
...
f += (direct + indirect) / (rr_prob)
return f
```



# First Attempt at Next Event Estimation

- Too bright! Better get some sunglasses to look at this...




BRDF Sampling



Current attempt



- Question: what happens if an indirect sample eventually hits a light?
- Indirect sample is accidentally direct, light is collected twice in same bounce!
- Due to the compensation of BRDF sampling, we end up with double the amount of light we should have!
- Idea: if we have double the amount of light, can we just divide by 2?




# Second Attempt at Next Event Estimation

```
function Li(v_inv, D)
...
f = x.emit
...
// Apply Russian Roulette at some point
...
direct = <direct lighting with light source sampling>
...
indirect = <indirect light (recursive) with BRDF sampling>
...
weight = (D == 0) ? 0.5 : 1 // halve indirect light after 1st bounce
f += weight * (direct + indirect) / (rr_prob)
return f
```




- The noise has significantly improved!
- Mixing several importance sampling techniques and weighting them...?
- It's multiple importance sampling!
- There are multiple ways to do MIS, let's quickly revisit some of them...



- Multiple MIS weightings to choose from:

- Averaging:  $w_i(x) = \frac{1}{N}$  for  $N$  techniques
- 1 or 0, depending on each new sample
- Balance heuristic (**Veach 1997**)

$$w_i(x) = \frac{p_i(x)}{\sum_{k=0}^N p_k(x)}$$



- Let's try something basic: assigning 0/1 weights to techniques
  - Assumption: light source sampling is better at **direct** light than BRDF
  - Keep BRDF sampling for indirect light, disable its direct light collection




# Removing Surface Emittance from Path Tracing

```
function Li(v_inv, D)
...
// f = x.emit
f = 0          // <-- 0 weight, removes direct light of BRDF sample at D-1!
...
// Apply Russian Roulette at some point
...
direct = <direct lighting with light source sampling>
...
indirect = <indirect light (recursive) with BRDF sampling> // unchanged
...
f += (direct + indirect) / (rr_prob) // 1 weight
return f
```



# Removing Surface Emittance from Path Tracing

- Looks better than averaging!
- But some information lost: light sources themselves don't seem to emit any light anymore...
- Logical, we removed emittance!
- It seems eliminating emittance altogether was too much...



- At the first bounce, there was no previous bounce for which we could compute the direct lighting with light source sampling
- I.e., we did not perform “next event estimation” at the 0<sup>th</sup> hit point, the camera (or viewpoint) itself
- Simple fix: actually, ignore emittance **most of the time**, except if the current hit point is the first hit after leaving the camera / eye
- Can use recursion depth to enable or disable emittance term




# Re-Enabling Emittance for Special Paths

```
function Li(v_inv, D)
    ...
    f = 0      // 0 weight, most of the time, except at first intersection
    if (D == 0)
        f = x.emit
    ...
    // Apply Russian Roulette at some point
    ...
    direct = <direct lighting with light source sampling>
    ...
    indirect = <indirect light (recursive) with BRDF sampling>
    ...
    f += (direct + indirect) / (rr_prob)
    return f
```




# Diffuse Next Event Estimation – Unlocked!



# More Diffuse Next Event Estimation Results

- BRDF importance sampling vs. next event estimation
  - In many cases, significant improvement of quality
  - Same number of samples, very similar runtime (NEE slightly slower)



# More Diffuse Next Event Estimation Results



# More Diffuse Next Event Estimation Results

- Next event estimation is highly effective, but still no silver bullet




- Always a more challenging scene to push your renderer to its limit...




# But is Using 0/1 Weights Really the Best Choice?

- Effectively, we now use light source sampling for all direct light
  - Using light source over BRDF sampling often improves direct lighting



White box, BRDF sampling



White box, light source sampling

- But, as usual, the benefit very much depends on the input scene




# Direct Light: Light Source Sampling is Not Always Better...

## ■ Scene with small and large area light: BRDF/light source sampling




BRDF sampling: large grey light works well,  
small blue light causes a lot of noise




Light source sampling: small blue light works  
well, large grey light causes a lot of noise



# How Sampling Methods Can Fail (Diffuse Materials)



Importance sampling the BRDF, we may miss the smaller light sources



Light source sampling a large, overhead light source (sky) concentrates more samples at flat angles – the opposite of importance sampling!

- Light source sampling can struggle even more with other BRDFs!
- But remember, there is another MIS option: the balance heuristic




# Direct Lighting with the Balance Heuristic


- Implementation in direct lighting integrator is not too complicated
  - Randomly choose technique and weight result using balance heuristic



BRDF sampling



Light source sampling




MIS, balance heuristic




- Consider both light sampling methods in every bounce
- Weight on first intersection must still be 1 for BRDF sampling
- Compute probabilities for selected sample with both techniques, use the balance heuristic to compute adequate MIS weights
- Tricky: different sampling methods are evaluated at different times
  - BRDF sample contribution is found when hitting an emitting surface
  - Contribution of light source sampling is calculated one bounce prior



# Path Tracing (NEE) with 0/1 Weights vs Balance Heuristic



- Soon, we will add some more exciting BRDFs for materials!
- Fully specular mirror materials are easy to simulate, however, they need extra care with NEE
- Naïve reflections can miss light!
- Why? Join us in the upcoming Materials lecture to find out...



- [1] Conquering noisy images in ray tracing with next event estimation:  
<https://developer.nvidia.com/blog/conquering-noisy-images-in-ray-tracing-with-next-event-estimation/>
- [2] Dittebrandt, A., Hanika, J., & Dachsbacher, C. (2020). Temporal Sample Reuse for Next Event Estimation and Path Guiding for Real-Time Path Tracing. In Eurographics Symposium on Rendering - DL-only Track. The Eurographics Association.
- [3] Guo, J., Eisemann, M., & Eisemann, E. (2020). Next Event Estimation++: Visibility Mapping for Efficient Light Transport Simulation. Computer Graphics Forum, 39(7), 205-217.
- [4] If you didn't like our explanations of next event estimation, perhaps Jacco Bikker can do a better job:  
<http://www.cs.uu.nl/docs/vakken/magr/2015-2016/slides/lecture%2008%20-%20variance%20reduction.pdf>

